

B-Trees (Chapter 18)

Bayer and McCreight (1972)



What does the “B” stand for?





Background

- › A **B-tree** is a balanced search tree which is used to store the keys to large file and database systems
- › Because:
 - only a part of a B-tree can be stored in primary (main) memory at a time and
 - secondary storage is much, much slower than primary memoryit is advantageous to read/write as much information as possible from/to secondary storage and thereby **minimize disk accesses**



- › Each node of a B-tree therefore is usually as large as a disk page and has a branching factor between 50 and 2000 depending on the size of the key relative to the size of a page
- › Hence, a B-tree is **wide** and **shallow**



Five properties of a B-tree

Property 1 *

Each internal node has
a minimum degree of t (except for the root) and a
maximum degree of $2t$.

* The notation B-tree ($t=k$) is used to denote a B-tree with a minimum degree of k



Five properties of a B-tree


Property 2

Each node has
a minimum of $t-1$ keys (except the root) and a
maximum of $2t-1$ keys.
If an internal node has degree d ($t \leq d \leq 2t$) then
it stores $d-1$ keys.




Properties 1 and 2 for B-tree ($t=3$)

Minimum degree



Maximum degree





Five properties of a B-tree

Property 3

The keys for each node are stored in ascending order.

Therefore,

$K_1 < K_2 < \dots < K_{d-1}$
for a node with degree d .



Five properties of a B-tree


Property 4

Key values in the i^{th} subtree of an internal node with degree d are

less than K_1 for $i=1$,
fall between K_{i-1} and K_i for $2 \leq i \leq d-1$, and
greater than K_{d-1} for $i=d$.



Properties 3 and 4 for B-Tree ($t=3$)





Five properties of a B-tree

Property 5

All leaf nodes have the **same** depth.



Maximum height h of a B-tree on n keys

Theorem (p489)

The maximum height on n keys for a B-tree with minimum degree t is

$$h \leq \log_t (n+1)/2 \text{ which is } O(\log n).$$




Exercises

- › Look up and review the proof on the previous slide.
- › Calculate the minimum height on n keys for a B-tree with minimum degree t .
- › Show all the legal B-trees ($t=2$) that represent the keys $\{1, 2, 3, 4, 5\}$.
- › For what values of t is the example tree on the next slide a legal B-tree?



Example tree *



* From CLRS (3rd Edition)



Data structure

```
class Node<T>
{
    private int n;           // number of keys
    private bool leaf;       // true if a leaf node; false otherwise
    private T[ ] key;        // array of keys
    private Node<T>[ ] c;    // array of child references

    public Node (int t) {
        n = 0;
        leaf = true;
        key = new T[2*t - 1];
        c = new Node<T>[2*t];
    }
}
```



Data structure

```
class BTree<T> where T : IComparable
{
    private Node<T> root;           // root node
    private int t;                  // minimum degree

    public BTree (int t) {
        this.t = t;
        root = new Node<T>(t);
    }

    ...
}
```




Primary methods

- › public bool Search (T k)
 - returns true if the key k is found in the B-tree; false otherwise
- › public void Insert (T k)
 - insert key k into the B-tree
 - duplicate keys are not permitted
- › public void Delete (T k)
 - delete key k from the B-tree



Support method

- › private void Split (Node<T> x, int i)
 - splits the i^{th} (full) child of x into 2 nodes



“One-pass” algorithms

- › The primary methods Search, Insert, and Delete are implemented as “one-pass” algorithms where each algorithm proceeds down the B-tree from the root **without** having to back up to a node *
- › This minimizes the number of read and write operations from secondary memory (disk)

* with one exception for deletion




Search

› Basic strategy

- Set p to the root node
- If the given key k is found at p then return true
- If the key k is not found and p is a leaf node then return false
else move to the subtree which would otherwise contain k




Search for key $k = 35$






Search for key $k = 7$






Search for key $k = 80$





Search for key $k = 20$





Time complexity

- › The length of the path from the root to a leaf node is $O(\log n)$
- › At each node along the path, it takes $O(t)$ time to determine which path to follow
- › Total time complexity is $O(t \log n)$



Insert


› Basic strategy

- Descend along the path based on the given key k from the root to a leaf node. Each node along the path though cannot be **full** i.e., have $2t-1$ keys. Therefore, **before** moving down to a node p (including the root), check first if p is full. If p is full then split p into two where the last $t-1$ keys of p are placed in a new node q and the median (middle) key of p is moved up to the parent node. Note: The parent node is guaranteed to have enough space to store the extra key (Why?)
- Insert the given key k at the leaf node. If the key is found along the path from the root to the leaf node, then no insertion takes place.




Insert V into a B-tree ($t=4$)

here




← not full

there



← full

Split method





← not full




Insert V into a B-tree ($t=4$) when the root is full

here

there




Height increases
by 1






In general

here






← not full

there




← full ($2t-1$ keys)




← not full




Initial B-Tree ($t=3$) \Rightarrow Do not descend to nodes with 5 keys

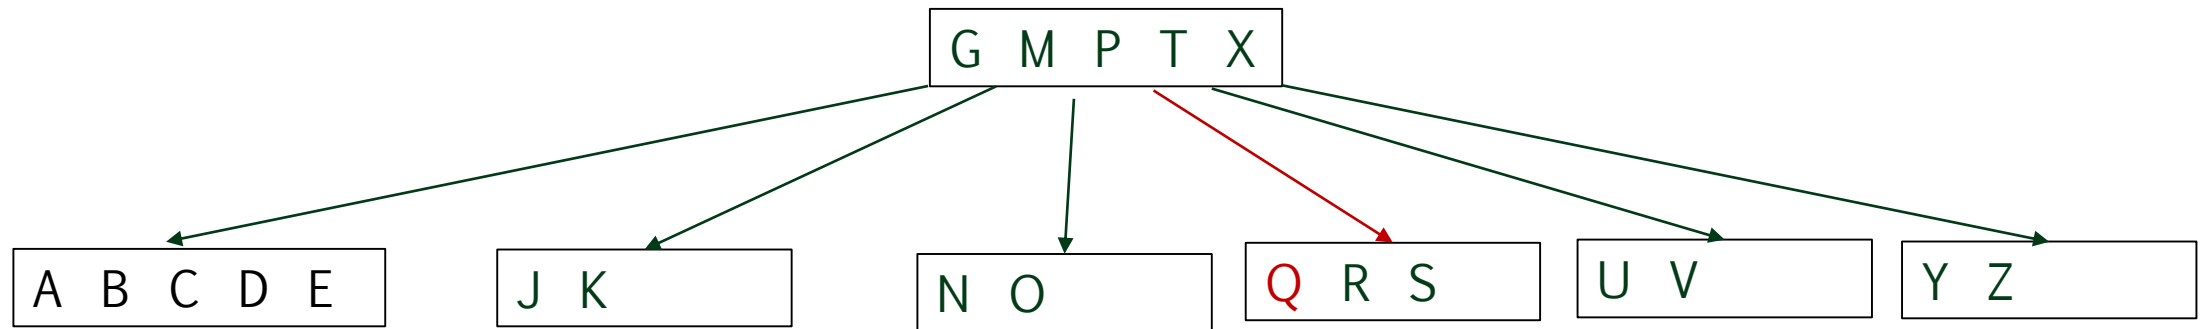



Insert B







Insert Q






Insert L





Insert F





Observations

- › Only descend to a node if it is **not full**
- › The B-tree **only** grows in height when a key is inserted into a B-tree with a full root i.e., when the root node is split
- › When a node is split into two, each of the two resultant nodes is assigned the minimum number of keys i.e., **$(t-1)$**



Time complexity

- › The length of the path from the root to a leaf node is $O(\log n)$
- › At each node along the path, it takes $O(t)$ time:
 - To determine which path to follow and
 - To split a node (if required)
- › Total time complexity is $O(t \log n)$



Exercises

- › Randomly insert a permutation of the alphabet into initially empty B-trees ($t=2$, $t=3$). Call the resultant trees B2 and B3.
- › Calculate the minimum and maximum heights of B-trees ($t=2$, $t=3$) for 26 keys.
- › Explain how to find the predecessor of a given key in a B-tree.



Delete

› Basic strategy

- Descend along the path based on the given key k from the root until the key is found. Each node along the path though must have at **least t keys**. Therefore, **before** moving down to a node p (excluding the root), check if p has at least t keys. If not, node p either:
 - › borrows a key from a sibling (if possible) or
 - › merges with an adjacent node where the $t-1$ keys of each node plus one from the parent node yield a single node with $2t-1$ keys. Note: The parent node is guaranteed to have an extra key (Why?)




› Basic strategy (cont'd)


- If the given key k is found at a leaf node, delete it. If the given the key k is found at an internal node p , then:
 - › if the child node q that precedes k has t keys, then recursively delete the predecessor k' of k in the subtree rooted at q and replace k with k' .
 - › if the child node r that succeeds k has t keys, then recursively delete the successor k' of k in the subtree rooted at r and replace k with k' .
 - › otherwise, merge q and r with k from the parent to yield a single node s with $2t-1$ keys. Recursively delete k from the subtree s .
- If key k is not found, no deletion takes place.




Initial B-Tree (t=3) => Do not descend to nodes with 2 keys or less





Delete F







Delete M



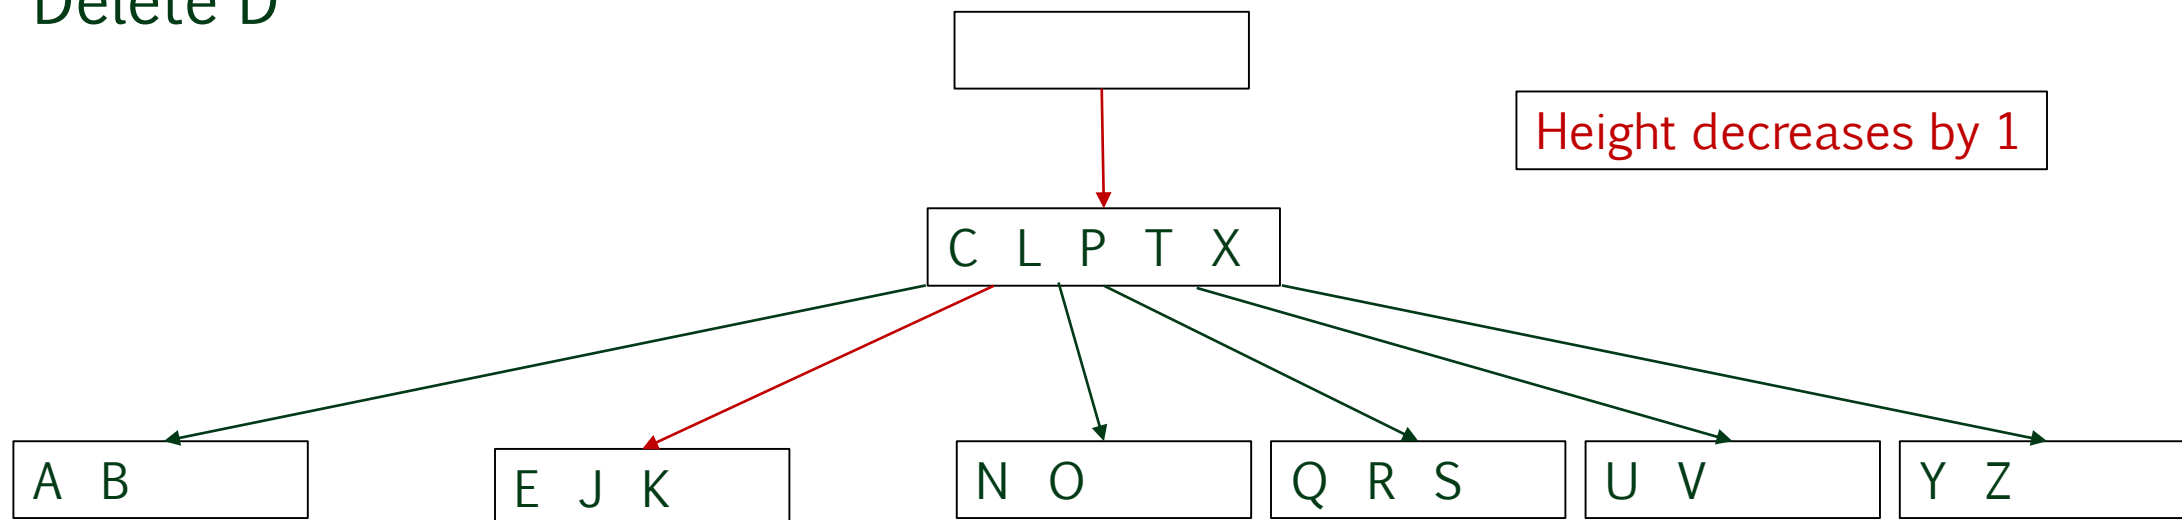


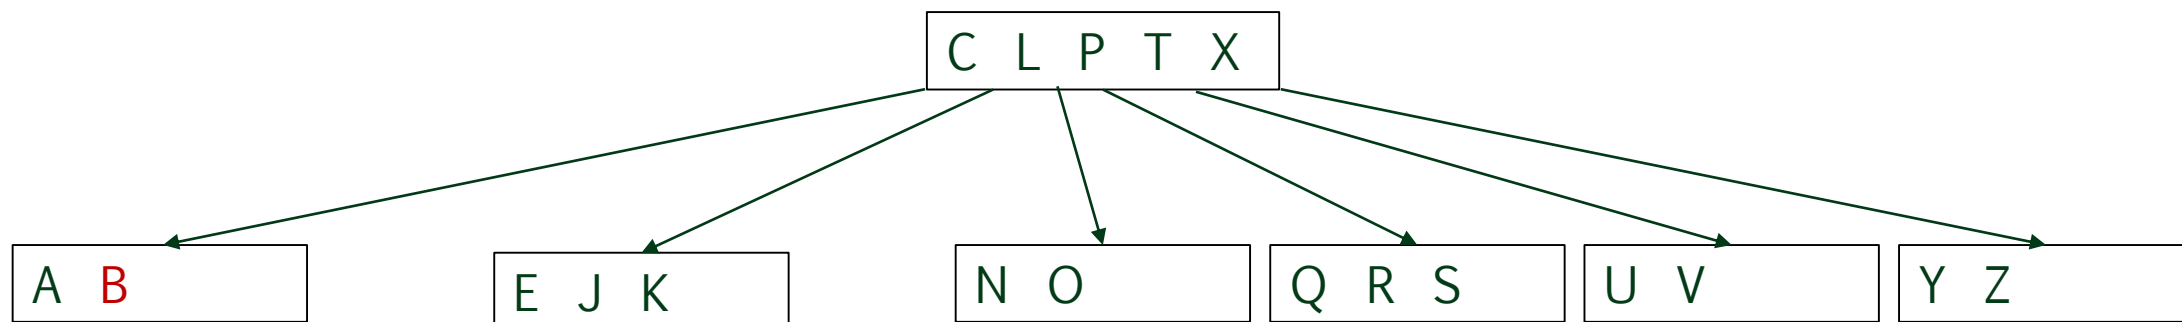
Delete G



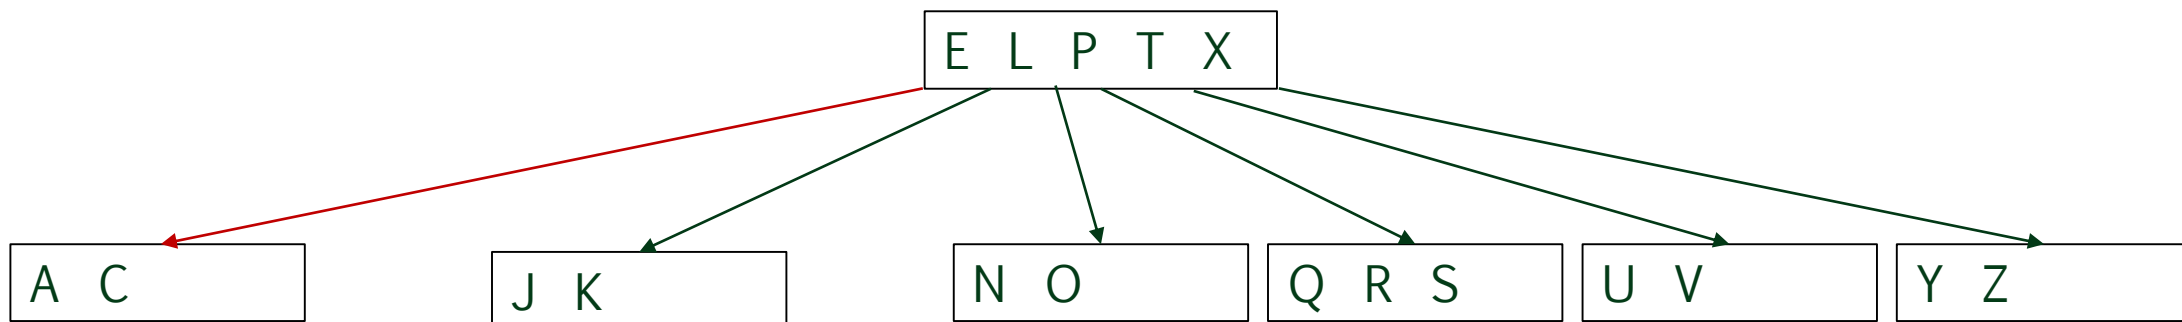


Delete D





Delete B





Observations

› The height of a B-tree **only** decreases if:

- the root node has a single key, and
- its two child nodes have $t-1$ keys each.

In this case, the two child nodes along with the key at the root are merged into one node with $2t-1$ keys. This node now serves as the new root of the B-tree.



Time complexity

- › Although the Delete method is quite complicated, its time complexity is the same as Search and Insert, $O(t \log n)$.



Exercises

- › Successively delete in the same order the keys that were inserted to construct B2 and B3 (from the last exercise).
- › Show why a node may need to be revisited (reread from disk) in the Delete method.
- › Show that the number of read and write disk accesses is $O(h)$ for Search, Insert, and Delete where h is the height of the B-tree.